

Twitter-like Scalable Backend System - Technical Documentation (Java Spring Boot)

1. System Architecture Overview

1.1 Objective

Design a scalable backend system similar to Twitter that can support millions of users with a focus on high availability, low latency, and modular architecture.

1.2 High-Level Architecture

Clients (Web, Mobile)



API Gateway (ALB / Spring Cloud Gateway)



Spring Boot Microservices (Fargate / EC2):

- Auth Service
- User Service
- Tweet Service
- Follow Service
- Timeline Service
- Notification Service

// Additional services:

- Search Service (Elasticsearch)

- Analytics Service

- Content Moderation Service

- Rate Limiting Service



Databases:

- PostgreSQL (Amazon RDS)
- Cassandra / DynamoDB



Event Queue:

- Apache Kafka (Amazon MSK)



Feed Workers (EC2 Auto Scaling Group)



Cache:

- Redis (Amazon ElastiCache)



Media:

- Amazon S3 / MinIO

1.3 Deployment Strategy

- Stateless services run on **AWS Fargate**
- Heavy background jobs and feed generation run on **EC2 Auto Scaling Group**
- Persistent storage via **RDS, ElastiCache, MSK, and S3**

2. Database Design

2.1 Relational Schema (PostgreSQL)

Users Table

```
CREATE TABLE users (  
  id BIGSERIAL PRIMARY KEY,  
  username VARCHAR(50) UNIQUE NOT NULL,  
  email VARCHAR(255) UNIQUE NOT NULL,  
  password_hash TEXT NOT NULL,  
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

Tweets Table

```
CREATE TABLE tweets (  
  id BIGSERIAL PRIMARY KEY,  
  user_id BIGINT REFERENCES users(id),  
  content TEXT NOT NULL,  
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

Follows Table

```
CREATE TABLE follows (  
  follower_id BIGINT REFERENCES users(id),  
  followee_id BIGINT REFERENCES users(id),  
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
  PRIMARY KEY(follower_id, followee_id)  
);
```

Likes Table

```
CREATE TABLE likes (  
  user_id BIGINT REFERENCES users(id),  
  tweet_id BIGINT REFERENCES tweets(id),  
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
  PRIMARY KEY(user_id, tweet_id)  
);
```

2.2 NoSQL Schema (Cassandra / DynamoDB)

User Feed Table

Table: user_feed
Partition Key: user_id
Sort Key: timestamp DESC
Attributes: tweet_id, tweet_content, author_id

3. API Design (Spring Boot + REST)

3.1 User APIs

Method	Endpoint	Description
POST	/auth/register	Register user
POST	/auth/login	Authenticate user
GET	/users/{id}	Get user profile
POST	/follow/{id}	Follow user
DELETE	/unfollow/{id}	Unfollow user

3.2 Tweet APIs

Method	Endpoint	Description
POST	/tweets	Create tweet
GET	/tweets/{id}	Get tweet
DELETE	/tweets/{id}	Delete tweet
POST	/tweets/{id}/like	Like a tweet
DELETE	/tweets/{id}/like	Unlike a tweet

3.3 Timeline APIs

Method	Endpoint	Description
GET	/timeline/home	Get home timeline
GET	/timeline/user/{id}	Get user's tweets

4. Core Features Implementation

4.1 User Registration and Authentication

- Spring Security + JWT + BCrypt
- Secure storage of password hashes
- Stateless authentication via tokens

4.2 Posting Tweets

- Tweets stored in PostgreSQL
- Kafka event emitted to timeline workers
- Worker pushes tweets to followers' feeds (fan-out)

4.3 Feed Generation

- **Hybrid Model:**
 - Fan-out on write for users with <10K followers
 - Fan-out on read for influencers/celebrities
- Timeline data cached in Redis (ElastiCache)

4.4 Likes and Follows

- Like and follow events trigger Kafka messages
- Notification service consumes events and inserts into user_notification table

4.5 Notification Service

- Kafka-based pub/sub
- Users notified of likes, mentions, follows
- Push to frontend via WebSocket or polling

4.6 Media Upload

- Spring Boot handles metadata
- Upload via pre-signed URL to Amazon S3

5. Caching Strategy

Data	Tool	TTL
Home Timeline	Redis	1-5 min
Tweets	Redis	10 min
User Profiles	Redis	30 min

6. Monitoring & Observability

- **AWS CloudWatch:** logs, metrics, alerts
- **AWS X-Ray:** distributed tracing
- **Prometheus + Grafana:** service metrics (optional)

7. Deployment & Scaling

7.1 Fargate Services

- Auth, Tweet, User, Follow, Notification APIs
- Scales automatically per demand

7.2 EC2 Services

- Feed generation workers

- Kafka consumers

7.3 Managed Services

- PostgreSQL (RDS)
- Kafka (MSK)
- Redis (ElastiCache)
- S3 (Media)

8. Conclusion

This Spring Boot-based backend is modular, scalable, and production-ready for a social media app like Twitter. It leverages AWS Fargate where serverless scaling is ideal, and EC2 where fine-tuned control is needed for high-load background processing.

The system follows modern architecture practices with event-driven design, hybrid feed modeling, and strong observability, ensuring it can support millions of users reliably and efficiently.